

# APMTH 207: Advanced Scientific Computing:

## Stochastic Methods for Data Analysis, Inference and Optimization

### Homework 11

Harvard University

Spring 2017

Instructors: Rahul Dave

Due Date: Monday, April 23rd, 2018 at 11:59pm

#### Instructions:

- Upload your final answers as an iPython notebook containing all work to Canvas.
- Structure your notebook and your work to maximize readability.

```
In [1]: import numpy as np
import pandas as pd
import time

import pymc3 as pm
import theano.tensor as T

import matplotlib
import matplotlib.pyplot as plt

import seaborn as sns
sns.set_style("whitegrid", {'axes.grid' : False})
sns.set_context('talk')
%matplotlib inline
```

## The AM207 Cambridge Nursery

A plant nursery in Cambridge is experimentally cross-breeding two types of hibiscus flowers: blue and pink. The goal is to create an exotic flower whose petals are pink with a ring of blue on each.

There are four types of child plant that can result from this cross-breeding:

- Type 1: blue petals
- Type 2: pink petals
- Type 3: purple petals
- Type 4: pink petals with a blue ring on each (the desired effect).

Out of 197 initial cross-breeding, the nursery obtained the following distribution over the four types of child plants:

$$Y = (y_1, y_2, y_3, y_4) = (125, 18, 20, 34)$$

where  $y_i$  represents the number of child plants that are of type  $i$ .

The nursery then consulted a famed Harvard plant geneticist, who informed them that the probability of obtaining each type of child plant in any single breeding experiment is as follows:

$$\frac{\theta + 2}{4}, \frac{1 - \theta}{4}, \frac{1 - \theta}{4}, \frac{\theta}{4}.$$

Unfortunately, the geneticist did not specify the quantity  $\theta$ .

Clearly, the nursery is interested in understanding how many cross-breeding they must perform, on average, in order to obtain a certain number of child plants with the exotic blue rings. To do this they must be able to compute  $\theta$ .

The owners of the nursery, being top students in AM207, decided to model the experiment in hopes of discovering  $\theta$  using the results from their 197 initial experiments.

They chose to model the observed data using a multinomial model and thus calculated the likelihood to be:

$$p(y|\theta) \propto (2 + \theta)^{y_1} (1 - \theta)^{y_2 + y_3} \theta^{y_4}$$

Being good Bayesians, they also imposed a prior on  $\theta$ ,  $\text{Beta}(a, b)$ .

Thus, the posterior is:

$$p(\theta|Y) = (2 + \theta)^{y_1} (1 - \theta)^{y_2 + y_3} \theta^{y_4} \theta^{a-1} (1 - \theta)^{b-1}.$$

If the nursery owners are able to sample from the posterior, they would be able to understand the distribution of  $\theta$  and make appropriate estimates.

## Problem 1. Sampling using data augmentation

Realizing that it would be difficult to sample from the posterior directly and after being repeatedly frustrated by attempts of Metropolis-Hastings and Gibbs sampling for this model, the nursery owners decided to augment their model and hopefully obtain a friendlier looking distribution that allows for easy sampling.

They augment the data with a new variable  $z$  such that:

$$z + (y_1 - z) = y_1.$$

That is, using  $z$ , we are breaking  $y_1$ , the number of type I child plants, into two subtypes. Let the probability of obtaining the two subtype be  $1/2$  and  $\theta/4$ , respectively. Now, we can interpret  $y_1$  to be the total number of trials in a binomial trial. Thus, the new likelihood can be written as

$$p(y, z|\theta) \propto \binom{y_1}{z} \left(\frac{1}{2}\right)^{y_1-z} \left(\frac{\theta}{4}\right)^z (1-\theta)^{y_2+y_3} \theta^{y_4}$$

Derive the joint posterior  $p(\theta, z|y)$  and sample from it using Gibbs sampling.

Visualize the distribution of theta and, from this distribution, estimate the probability of obtaining a type 4 child plant (with the blue rings) in any cross-breeding experiment.

## Answer to Problem 1

We know  $\theta \sim \text{Beta}(a, b)$ , and  $p(y, z|\theta) \propto \binom{y_1}{z} \left(\frac{1}{2}\right)^{y_1-z} \left(\frac{\theta}{4}\right)^z (1-\theta)^{y_2+y_3} \theta^{y_4}$ . Thus

$$p(\theta, z|y) \propto p(y, z|\theta) p(\theta) = \binom{y_1}{z} \left(\frac{1}{2}\right)^{y_1-z} \left(\frac{\theta}{4}\right)^z (1-\theta)^{y_2+y_3} \theta^{y_4} \theta^{a-1} (1-\theta)^{b-1}$$

To obtain conditionals, we keep the relevant variables.

$$p(\theta|y, z) \propto \theta^{z+y_4+a-1} (1-\theta)^{y_2+y_3+b-1}$$

i.e.,

$$\theta|y, z \sim \text{Beta}(z + y_4 + a, y_2 + y_3 + b).$$

And

$$p(z|y, \theta) \propto \binom{y_1}{z} \left(\frac{2}{\theta+2}\right)^{y_1-z} \left(\frac{\theta}{\theta+2}\right)^z$$

i.e.,

$$z|y, \theta \sim \text{Binomial}(y_1, \frac{\theta}{\theta+2}).$$

We can implement Gibbs sampling using the above conditionals. And we simply choose a uniform prior for  $\theta$  (i.e.,  $a = b = 1$ ).

```
In [2]: def corrplot(trace, maxlags=50):  
        plt.acorr(trace-np.mean(trace), normed=True, maxlags=maxlags);  
        plt.xlim([0, maxlags])  
  
        def effective_sample_size(data, step=1):
```

```

# References:
# https://code.google.com/p/biopy/source/browse/trunk/biopy/bayesianStats.py?r=67
# https://am207.github.io/2018spring/wiki/tetchygibbs.html

n = len(data)
assert n > 1
maxlags = min(n//3, 1000)

gamma_stat = [0, ] * maxlags

var_stat = 0.0

if type(data) != np.ndarray:
    data = np.array(data)

data_normed = data - data.mean()

for lag in range(maxlags):
    v1 = data_normed[:n-lag]
    v2 = data_normed[lag:]
    v = v1 * v2
    gamma_stat[lag] = sum(v) / len(v)

    if lag == 0:
        var_stat = gamma_stat[0]
    elif lag % 2 == 0:
        s = gamma_stat[lag-1] + gamma_stat[lag]
        if s > 0:
            var_stat += 2 * s
        else:
            break

act = step * var_stat / gamma_stat[0]
ess = step * n / act

return ess

def print_ess(data):
    ess1 = effective_sample_size(data[:, 0])
    ess2 = effective_sample_size(data[:, 1])
    print('Effective size for theta:', ess1, ' of', len(data), 'samples; effective rate:', ess1/len(data))
    print('Effective size for z:', ess2, ' of', len(data), 'samples; effective rate:', ess2/len(data))

class Gibbs:
    def __init__(self, a=1, b=1, y=np.array([125, 18, 20, 34])):
        self.a = a
        self.b = b
        self.y = y

    def run(self, n=20000, x_init=np.array([0.5, 0]), seed=0):
        a = self.a
        b = self.b
        y = self.y

```

```
a0 = y[3] + a
b1 = y[1] + y[2] + b

samples = np.empty((n+1, 2))
samples[0] = x_init

if seed is not None:
    np.random.seed(seed)

for i in range(1, n+1):
    a1 = samples[i-1, 1] + a0
    samples[i, 0] = np.random.beta(a1, b1)

    p = samples[i, 0] / (samples[i, 0] + 2)
    samples[i, 1] = np.random.binomial(y[0], p)

self.n = n
self.samples = samples[1:]
return self

def process(self, burnin=0, thin=1):
    self.burnin = burnin
    self.thin = thin
    self.samples2 = self.samples[burnin::thin]
    return self
```

```
In [3]: g0 = Gibbs().run().process(thin=2)
```

```
In [4]: plt.figure(figsize=(10, 3))

plt.subplot(1, 2, 1)
corrplot(g0.samples[:, 0])
plt.ylabel(r'autocorrelation($\theta$)')


plt.subplot(1, 2, 2)
corrplot(g0.samples[:, 1])
plt.ylabel('autocorrelation($z$)')
plt.tight_layout()

print('No burnin and thining:')
print_ess(g0.samples)
```

No burnin and thining:

Effective size for theta: 14612.3871292 of 20000 samples; effective rate: 0.73061935646

Effective size for z: 14841.4240569 of 20000 samples; effective rate: 0.742071202847




```
In [5]: plt.figure(figsize=(10, 3))

plt.subplot(1, 2, 1)
corrplot(g0.samples2[:, 0])
plt.ylabel(r'autocorrelation($\theta$)')

plt.subplot(1, 2, 2)
corrplot(g0.samples2[:, 1])
plt.ylabel('autocorrelation($z$)')
plt.tight_layout()

print('burnin = {}, thinning = {}'.format(g0.burnin, g0.thin))
print_ess(g0.samples2)


burnin = 0, thinning = 2:
Effective size for theta: 9592.65625146 of 10000 samples; effective rate: 0.959265625146
Effective size for z: 9406.20749842 of 10000 samples; effective rate: 0.940620749842
```



As we can see, autocorrelations are negligible after thinning at 2.

```
In [6]: thetas = g0.samples2[:, 0]

plt.plot(thetas);
plt.ylabel(r'$\theta$');
```






We can visualize the distribution of  $\theta$  as follows.

```
In [7]: sns.distplot(thetas);
plt.xlabel(r'$\theta$');

print('mean(theta) = {}'.format(np.mean(thetas)))
print('std(theta) = {}'.format(np.std(thetas)))
```

```
mean(theta) = 0.6229266753714496
std(theta) = 0.05064704358302179
```



```
In [8]: print('The estimate of the probability of obtaining a type 4 child plant is
{:.4f}'.format(np.mean(thetas/4)))
print('The 2-sigma envelope is [{:.4f}, {:.4f}].'.format(np.mean(thetas/4)-2*
p.std(thetas/4), \
np.mean(thetas/4)+2*
p.std(thetas/4)))
```

```
The estimate of the probability of obtaining a type 4 child plant is 0.1557.
The 2-sigma envelope is [0.1304, 0.1811].
```

## Problem 2. Finding the MLE using Expectation Maximization

Treat the augmented model as a latent variable model.

### Part A.

Write down an expression (up to unimportant constants - you must decide what unimportant means) for each of the following:

- (1) the observed data log likelihood
- (2) the complete(full) data log likelihood

**Hint:** You should already have the observed data likelihood and the complete data likelihood from Problem 1, you just need to take their logs for this problem.

- (3) the Auxiliary function,  $Q(\theta, \theta^{(t-1)})$ , or the expected complete(full) data log likelihood, defined by
 
$$Q(\theta, \theta^{(t-1)}) = \mathbb{E}_{Z|Y=y, \Theta=\theta^{(t-1)}} [\text{the complete data log likelihood}]$$

In other words  $Z|Y=y, \Theta=\theta^{(t-1)}$  is  $q(z, \theta_{old})$  from lecture at the end of the E-step and  $Q$  is the z-posterior expectation (at  $\theta_{old}$ ) of the full data log likelihood, which is the ELBO minus the entropy of  $q$  (which being evaluated at  $\theta_{old}$  is not dependent on  $\theta$  and thus irrelevant for maximization).

### Part B:

We will maximize the likelihood through Expectation Maximization (EM). In order to perform EM, we must iterate through the following steps

- (Expectation) Compute the Auxiliary function,  $Q(\theta, \theta^{(t-1)})$  (the expectation of the full data likelihood)
- (Maximization) Compute  $\theta^t = \operatorname{argmax}_{\theta} Q(\theta, \theta^{(t-1)})$

Thus, you must compute exact formulae for the following:

1. the Auxiliary function,  $Q(\theta, \theta^{(t-1)})$ , for a given  $\theta^{(t-1)}$ . That is, compute the expectation of the complete data log likelihood.
2.  $\theta^t$ , by maximizing the Auxiliary function  $Q(\theta, \theta^{(t-1)})$ .

**Hint:** You don't actually need to do any difficult optimization for the M-step. After taking the expectation of the complete data log likelihood in the E-step, match your  $Q(\theta, \theta^{(t-1)})$  to the log pdf of a familiar distribution, then use the known formula for the mode of this distribution to optimize  $Q(\theta, \theta^{(t-1)})$ .

Use these to **estimate the MLE** of  $\theta$  using EM (choose your own reasonable criterion for convergence).

## Answer to Problem 2 Part A

Let  $n$  be the total number of initial cross-breeding ( $n = 197$  in this case); let  $C_i$  represent unimportant constants.

### (1) the observed data log likelihood

$$y|\theta \sim \text{Multinomial}\left(y; n, \frac{\theta + 2}{4}, \frac{1 - \theta}{4}, \frac{1 - \theta}{4}, \frac{\theta}{4}\right)$$

$$\log p(y|\theta) = y_1 \log(2 + \theta) + (y_2 + y_3) \log(1 - \theta) + y_4 \log \theta + C_0$$

### (2) the complete (full) data log likelihood

$$\log p(y, z|\theta) = \log \binom{y_1}{z} + (y_1 - z) \log \frac{1}{2} + z \log \frac{\theta}{4} + (y_2 + y_3) \log(1 - \theta) + y_4 \log \theta + C_1$$

### (3) the Auxiliary function

We know from problem 1 that

$$z|y, \theta \sim \text{Binomial}(y_1, \frac{\theta}{\theta + 2})$$

. Thus,

$$\mathbb{E}_{z|y, \theta^{t-1}}(z) = \frac{y_1 \theta^{t-1}}{\theta^{t-1} + 2}.$$

Let  $F(z, \theta^{t-1})$  represent terms not involving  $\theta$  (including unimportant constants), then

$$\begin{aligned} Q(\theta, \theta^{t-1}) &= \mathbb{E}_{z|y, \theta^{t-1}}[\log p(y, z|\theta)] = \mathbb{E}_{z|y, \theta^{t-1}} \left[ \log \binom{y_1}{z} + (y_1 - z) \log \frac{1}{2} + z \log \frac{\theta}{4} + (y_2 + y_3) \log(1 - \theta) \right] \\ &= F(z, \theta^{t-1}) + \mathbb{E}_{z|y, \theta^{t-1}}[z \log \theta + y_4 \log \theta + (y_2 + y_3) \log(1 - \theta)] \\ &= F(z, \theta^{t-1}) + (\mathbb{E}_{z|y, \theta^{t-1}}(z) + y_4) \log \theta + (y_2 + y_3) \log(1 - \theta), \end{aligned}$$

where

$$\mathbb{E}_{z|y, \theta^{t-1}}(z) = \frac{y_1 \theta^{t-1}}{\theta^{t-1} + 2}.$$

## Answer to Problem 2 Part B

At E-step, we need to evaluate

$$\mathbb{E}_{z|y, \theta^{t-1}}(z) = \frac{y_1 \theta^{t-1}}{\theta^{t-1} + 2}.$$

At M-step, we know

$$Q(\theta, \theta^{t-1}) = F(z, \theta^{t-1}) + (\mathbb{E}_{z|y, \theta^{t-1}}(z) + y_4) \log \theta + (y_2 + y_3) \log(1 - \theta).$$

Let  $\frac{\partial Q(\theta, \theta^{t-1})}{\partial \theta} = 0$ , we get

$$\theta^t = \frac{y_4 + \mathbb{E}_{z|y, \theta^{t-1}}(z)}{y_2 + y_3 + y_4 + \mathbb{E}_{z|y, \theta^{t-1}}(z)}.$$

We can iterate the process until  $\theta$  converges.


```
In [2]: class EM:
    def __init__(self, y=np.array([125, 18, 20, 34])):
        self.y = y

    def run(self, theta_init=0, max_iter=1000, thres=1e-10):
        y = self.y
        zs = []
        thetas = [theta_init]
        for _ in range(max_iter):
            zs.append(y[0]*thetas[-1]/(thetas[-1]+2))
            thetas.append((y[3]+zs[-1])/(y[1]+y[2]+y[3]+zs[-1]))
            if abs(thetas[-1]-thetas[-2]) < thres:
                break
        self.zs = zs
        self.thetas = thetas
        return self
```

We can run the algorithm from several different initial  $\theta$ s to see whether they converge to the same point.

```
In [3]: ems = [EM().run(theta_init=t) for t in [0, 0.2, 0.4, 0.6, 0.8, 1]]

for em in ems:
    plt.plot(range(len(em.thetas)), em.thetas, '-.', alpha=0.7, label=r'$\theta_0 = {}'.format(em.thetas[0]));
plt.legend(bbox_to_anchor=(1, 1));
plt.xlabel('iteration');
plt.ylabel(r'$\theta$');
```



As we can see, they all converge to the same point.

```
In [4]: print('The estimate of the MLE of theta using EM is {:.4f}'.format(ems[0].thetas[-1]))
```

The estimate of the MLE of theta using EM is 0.6268.